

# Query Optimization

Ilaria Battiston \*

Winter Semester 2020-2021

---

\*All notes are collected with the aid of material provided by T. Neumann. All images have been retrieved by slides present on the TUM Course Webpage.

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Query Processing . . . . .	3
1.2	Query Optimization . . . . .	4
1.3	Query Execution . . . . .	5
<b>2</b>	<b>Textbook Query Optimization</b>	<b>6</b>
2.1	Algebra and tuples . . . . .	6
2.2	Canonical Query Translation . . . . .	7
2.3	Logical Query Optimization . . . . .	7
2.4	Physical Query Optimization . . . . .	8
<b>3</b>	<b>Join Ordering</b>	<b>8</b>
3.1	Greedy heuristics . . . . .	11
3.2	IKKBZ . . . . .	12
3.3	Maximum Value Precedence . . . . .	15
3.4	Dynamic Programming . . . . .	17

# 1 Introduction

Queries involve multiple steps to be performed. All techniques obviously depend on the physical characteristic of hardware, however there are algorithms and data structures able to help making them faster.

Query optimization is specifically important since SQL-like languages are declarative, hence do not specify the exact computation. Furthermore, the same query can be performed in multiple ways, and choosing the correct one is not trivial (depending on size, indices and more).

For instance, looking at cardinality of joins can either cause a Cartesian product of a hundred or several thousands, depending on the order in which they are executed. Often the human-written sequential order of instructions is not the most efficient one.

## 1.1 Query Processing

Query processing consists in:

1. Taking a text query as input;
2. Compiling and optimizing;
3. Extract the execution plan;
4. Executing the query.

Roughly, it can be differentiated between compile time and runtime system. Most systems strongly separate those phases, for example using so-called prepared queries.

```
SELECT S.NAME  
FROM STUDENTS S  
WHERE S.ID = ?
```

Compilation takes time and limits the number of queries which can be ran in parallel. This kind of query can be executed over and over providing a value `Q1(123)` without the compilation overhead.

Embedded SQL is another technique to optimize compilation time in the case of large periods of time between the two phases. The programming language compiler takes care of the SQL part as well, optimizing it.

Specifically, the steps executed in compile time are:

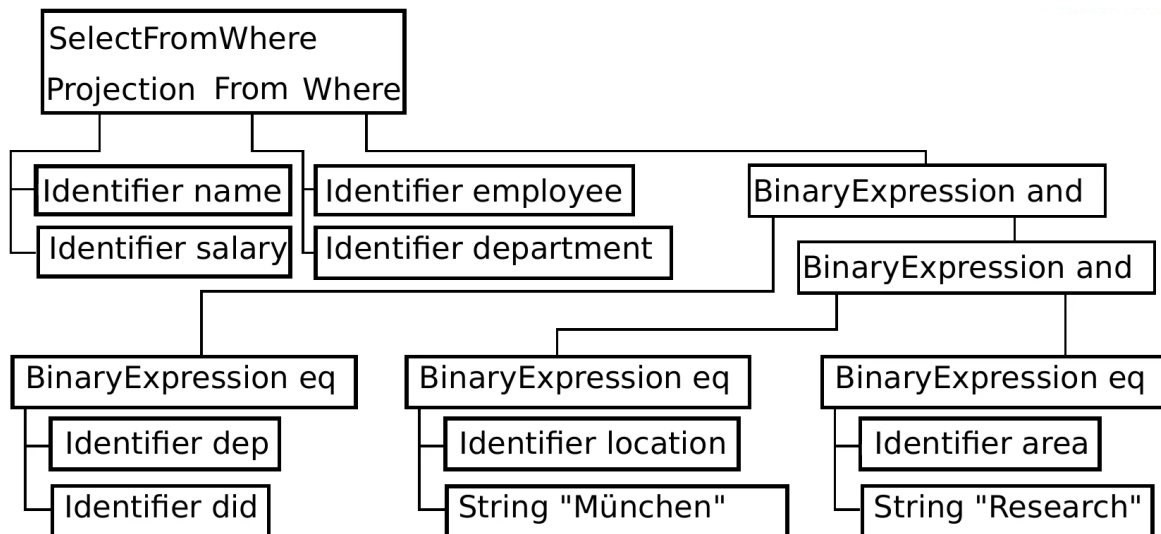
1. Parsing, AST production (abstract syntax to understand the structure);
2. Schema lookup, variable binding, type inference (semantic analysis about relations and columns, syntax check);
3. Normalization, factorization (bringing the query in abstract form, avoiding computing the same thing twice, evaluating expressions);
4. Unnesting, deriving predicates, resolution of views (the plan generator can finally construct a cost-based model);
5. Construction of execution plan;
6. Review, pushing joins and refining the plan in general;

7. Production of imperative plan (code generation).

Rewrite I involves steps 1-3, while 4 and 5 compose rewrite II.

Example (with views):

```
SELECT name, salary
FROM employee, department
WHERE dep = did
AND location = "Munich"
AND area = "Research"
```



Finally, the execution tree is built and polished, with the join operation on top and selects on nodes, reducing the amount of tuples to be joined. In other cases, such as regular expressions which are hard to evaluate, filtering can be done later.

The executable plan is a set of tuples containing variables or constants along and their type, with the main function loading query parameters, operations and allocated resources.

Usually query planners are much more complicated and have practical difficulties: for instance, a long list of AND/OR predicates (machine-generated in the order of hundreds of thousands) can make the binary tree recursion crash due to insufficient space in the stack.

## 1.2 Query Optimization

Possible goals of query optimization include minimizing response time, resource consumption, time to first tuple (producing the first tuple as quick as possible, for instance with search results) or maximizing throughput. This can be expressed as a cost function: most systems aim to minimize response time, having resources as constraints.

Algebraic optimization is a branch using relational algebra to find the cheapest expression equivalent to the original. However, finding the cheapest is a practically impossible problem: it is hard to test

for equivalence (numerical overflow, undecidable), the set of expressions is potentially huge and some algorithms are NP-hard (actual search space is limited and smaller than the potential one).

There are ways to transform numerical expressions in algebraic ones, yet they might be expensive: calculus is faster to evaluate than algebra.

Optimization approaches can be:

- Transformative, taking an algebraic expression and iteratively making small changes, not efficient in practice;
- Constructing, starting from small expressions and joining them, obtaining larger sets, usually the preferred approach.

### 1.3 Query Execution

Query execution is the last step, the one directly benefiting from optimization. In reality, operators can perform extremely specialized operations, treating data as bags (sets with duplicates) or streams.

#### 1.3.1 Relational and physical algebra

Relational algebra for query optimization includes the operators of projection, selection and join, plus some additional ones such as mapping and grouping.

However, this does not imply an implementation, which can have a great impact, hence less abstract operators are needed due to stream nature of data (and not sets):

- Sort, according to criteria;
- Temp, which materializes the input stream making further reads cheaper and releasing memory afterwards;
- Ship, sends the input stream to a different host (for distributed computing).

Furthermore, there are different kind of joins having different characteristics:

- Nested Loop Join, the textbook implementation which is very slow (quadratic) but supports all predicates;
- Blockwise Nested Loop Join, reading chunks of the first column in memory and then the second one once (requires memory, can be improved using hashing). The left side is read  $\frac{|L|}{|M|}$  times, which asymptotically is the same as above but in practice much faster;
- Sort Merge Join, scanning each column only once, but requiring sorted input ( $O(n \log n)$  plus linear runtime assuming no duplicates on at least one side);
- Hybrid Hash Join, partitioning columns and joining them in memory (linear time, but working only for equi-joins).

Hash Join is usually the fastest, but with very large amount of data sorting might be more efficient. Nested Loop Join works well when one of the sides has only one tuple.

Other operators also have different implementations, such as aggregation sorting and hashing, whose speed depends on the algorithm (quick sort, heap sort etc.) and can rely or not on memory.

Physical algebra is hence the output of the query optimizer, since it fixes the exact execution runtime and does not stop to the general approach like relational algebra. It mainly helps to select which operator to use before running the query, so that optimization leads to more efficient plans.

## 2 Textbook Query Optimization

Textbook query optimization involves techniques to perform a rough first optimization, which however is quite simple. There is a series of steps to translate raw SQL into logical and physical plans, each of them transforming input in a more optimal form.

The output is going to be executable, but still to be improved by non-trivial methods.

### 2.1 Algebra and tuples

Plain relational algebra is not sufficient itself: it needs to be revisited ensuring correctness (producing the same result) within a formal model. The most relevant problem to tackle is deciding whether two algebraic expressions are the same, but this is difficult in practice.

For instance, performing a selection before a join might be correct (and faster) in case the considered criterion is equality, but can give a different result than selecting after an outer join.

To remedy this issue, it is possible to guarantee that two expressions are equivalent, not accepting false positives yet allowing false negatives.

A formal definition of tuple is an unordered mapping from attribute names to values of a domain. A schema consists in a set of attributes with domain  $A(t)$ .

Tuple operations are:

- Concatenation, attaching one tuple to another regardless of ordering (union);
- Projection, producing a notation  $t.a$  in which it is possible to access single values or multiple  $t_{|\{a,b\}}$ , getting a subset of the schema.

A set of tuples with the same schema forms a relation. Sets naturally do not comply with real data, since they not allow duplicates, but are used for simplicity.

In most cases, sets and bags can be used interchangeably, but the optimizer considers different semantics: logical algebra operates on bags, physical algebra on streams and sets are only considered after an explicit duplicate elimination.

Set operations are the classic ones of union, intersection and difference, yet are subject to schema constraints. On bags, operations are performed on frequencies.

There are also free variables, which first must be bounded to be evaluated: they are essentials for predicates and algebra expressions, such as dependent joins.

It is important to note that projection removes duplicates within sets, while keeping them in bags.

There are equivalences for selection and projection useful to derive whether a different ordering produces the same output. For instance, applying selection twice is the same as applying it once with two criteria plus an AND. Commutative property also holds.

## 2.2 Canonical Query Translation

The canonical query translation transforms SQL into algebra expressions. The first approach involves some restrictions: it assumes no duplicates without aggregation and set operations.

The first step is translating the FROM clause:

$$F = \begin{cases} R_1 & k = 1 \\ ((\dots (R_1 \times R_2) \times \dots) \times R_k) & \text{else} \end{cases}$$

In short all relations are joined through a cross product. The next step is translating the WHERE clause:

$$W = \begin{cases} F & \text{there is no WHERE clause} \\ \sigma_p(F) & \text{otherwise} \end{cases}$$

The SELECT clause is translated starting from the projection  $a_1, \dots, a_n$  or \*. The expression is constructed:

$$S = \begin{cases} W & \text{if the projection is ALL} \\ \Pi_{a_1, \dots, a_n}(W) & \text{otherwise} \end{cases}$$

GROUP BY can also be translated, even though it is not part of the canonical translation. Let  $g_1, \dots, g_n$  be the attributes in the clause and *agg* the aggregations within SELECT:

$$G = \begin{cases} W & \text{there is no GROUP BY clause} \\ \Gamma_{g_1, \dots, g_m:agg}(W) & \text{otherwise} \end{cases}$$

HAVING is basically the same as WHERE, with the filter predicate on top of  $G$ .

## 2.3 Logical Query Optimization

Once obtained the relational algebra, equivalences span the potential search space and new expressions are derived thanks to them. Of course equivalence can be applied in both ways, hence it is relevant to decide which one works better, and conditions have to be checked as well. This, however, makes the search more expensive since there are plenty of alternatives.

To speed the process up, sometimes some equivalences are ignored, even the simplest ones (for instance when choosing the join algorithm).

Query plans can only be compared if there is a cost function, often needing details which are not available merely through relational algebra (what kind of join is being used): logical query optimization is still a heuristic and requires additional steps, since it is not enough to determine the runtime.

Most algorithms, therefore, use the following strategy:

- Organization of equivalences into groups;
- Directing equivalences, deciding the preferred side and rewriting rules to apply them sequentially to the initial expression, trying to reduce the size of intermediate results.

For example, a projection on the output of a join can be preferred to a join of a projection. It is important to keep in mind that tuples are being removed in the process, and this only applies in certain circumstances (regular expressions, high selectivity of join).

The rule of thumb is simply to eliminate the most tuples during the intermediate step, to then perform computationally expensive operation with the smallest amount of data.

To summarize, the phases are:

- Breaking up conjunctive selection predicates, since simpler predicates can be moved around easier;
- Pushing selections down, reducing the number of tuples early;
- Introducing joins, which are cheaper than cross product (linear time);
- Determining join order;
- Introducing and pushing down projections, removing redundant attributes.

Some SQL queries have limitations: selections sometimes cannot be pushed down, since there might be no join predicate between tables. Choosing a different join order allows further push down.

## 2.4 Physical Query Optimization

Physical query optimization adds execution information to the plan, allowing actual cost calculation and optimizing over data structures, access path and operator implementation.

Data may be sorted or materialized, introducing results which can be reused and deciding where to store them.

First of all, the access path is selected: lookup can be done through index or table scan, depending on the selectivity (fraction of the data satisfying the clause): in general, above 10% a table scan is recommended.

Scanning a table might be efficient since tuples are stored adjacent in memory; using index, instead, involves traversing a tree multiple times starting from the root.

Sometimes it is useful to just store in cache the output of a view, but that also depends on the query plan: intermediate results should actually be reused.

Operator selection is replacing a logical operator with a physical one, according to semantic restrictions (most operators require equi-join).

A blockwise nested loop join is generally better than a natural join; sort merge join and hash join are better than both. In general, hash join is the best if not reusing sorts. This process must be performed for all operators: sort join requires ordered tuples, distributed databases need local data and there are multiple ways to model the properties (hashing).

Sort merge join might outperform the hash join if the amount of data is much larger than the available memory.

Materializing, on the other side, is quite relevant for nested loop joins: the first pass is expensive, but the afterwards ones are way cheaper, making it essential for multiple consumers.

## 3 Join Ordering

Join ordering focuses on conjunctive queries with simple predicates of the type  $a_1 = a_2$  where the latter can be either an attribute or a constant (commonly algorithms assume join between attributes).




Relations may include selections or complex building blocks, however for simplicity filtering is ignored; having operators other than equality might cause differences within the query planner.

Ordering basically means choosing which relation to be joined first, placing entities in a graph and adding an edge whenever a predicate from a node is joined to another.

This kind of schema is defined as a query graph, in which edges consist in predicates and self loops represent equality with a constant. Usually cycles are pushed down, since algorithms only assume attributes.

Based on the query graph it is possible to obtain an overview of the complexity of the problem: there are different shapes which are treated differently.



1. Chains are the simplest kind of query, fairly common in practice;
2. Cycles (cyclic) are a chain with a closing edge, the easiest example of cycles;
3. Stars are mostly used in data warehouse, in which the center table has large dimension and the ones outside are relatively small, quite different to solve;
4. Cliques are instances in which every relation is joined with all the others, and are the hardest to optimize causing the worst runtime;
5. Trees are acyclic queries even if the level of nesting can be high;
6. Grids are also fairly hard and interesting for research.

Joins are represented with join trees, binary trees with operators as inner nodes and relations as leaves. The most common type is unordered (not distinguish left from right) without cross product, however algorithms might produce other variants.

There furthermore are different kinds of trees:

- Left-deep tree, in which joins only happen on the left side, easy to represent and implement through hash tables ( $n!$  trees with cross products);
- Right-deep tree ( $n!$ );
- Zig-zag tree, a combination of the previous ( $n!2^{n-2}$ );

- Bushy tree, a full binary tree (non-linear, harder to find optimal solutions but can be the most efficient in some cases,  $n!C(n-1) = \frac{(2n-2)!}{(n-1)!}$  where  $C$  represents a Catalan number).

It is relevant to notice that the number of leaf combinations and unlabeled trees grows exponentially, and increases even more with a flexible structure. However, nodes can often be swapped from left to right.

Another important information about joins is their selectivity:

$$f_{i,j} = \frac{|R_i \bowtie_{p_{i,j}} R_j|}{|R_i \times R_j|}$$

This depends on whether the attributes are a key, and gives an estimation of the result cardinality with the aid of assumptions and statistics.

Given a join tree, the cardinality can be computed recursively as the product of the selectivity function multiplied by the size of both relations. This allows easy calculations only requiring base cardinalities and independence of predicates:

$$C_{out}(T) = \begin{cases} 0 & T \text{ is a leaf} \\ |T| + C_{out}(T_1) + C_{out}(T_2) & T = T_1 \bowtie T_2 \end{cases}$$

This formula sums up the sizes of intermediate results, which are the ones causing more works. There are basic specific cost functions for joins, to be summed to the cost of single relations.

Algorithms are mainly designed for left-deep trees, and some of the cost functions do not work in practice, for instance in the case of cross products. Therefore, those indicators are mainly theoretical and work under strict assumptions. However, join ordering is a main factor regardless of the chosen cost methods.

A cost function is called symmetric if  $C_{impl}(e_1 \bowtie^{impl} e_2) = C_{impl}(e_2 \bowtie^{impl} e_1)$ . Commutativity can be ignored.

Most of the time, algorithms for query optimization tend to avoid cross products, despite the enormous number of possibilities to build a join tree: the only exception regards small relations.

### 3.0.1 Chains

Chains usually originate a left-deep tree: leaves can be ordered according to different degrees of freedom, as long as all the relations are joined.

The number of possible left-deep trees can be defined recursively:

$$\begin{cases} f(0) = 0 \\ f(1) = 1 \\ f(n) = 1 + \sum_{k=1}^{n-1} f(k-1) \cdot (n-k) \end{cases}$$

Adding  $R_n$  to all possible join trees can be done at any position following  $R_{n-1}$ . There are  $n-k$  join trees for  $R_n$ , plus one assuming it can be placed before  $R_{n-1}$  in the case of  $k=1$ . For  $R_{n-1}$  to be at  $k$ ,  $R_{n-k} - \dots R_{n-2}$  must be below it.

Solving the recurrence gives the closed form  $f(n) = 2^{n-1}$ , still exponential yet much less than the case with cross products.

A generalization to zig-zag can be made expecting the same result.

Bushy trees, on the other hand, are not so easy to obtain: each subtree must contain a subchain to avoid cross products, hence single relations should not be added. It is possible to create a whole chain  $R_1 - \dots - R_n$ , cut it and place it under another subtree, always considering commutativity.

This gives the formula:

$$f(n) = \begin{cases} 1 & n < 2 \\ \sum_{k=1}^{n-1} 2f(k) \cdot f(n-k) & n \geq 2 \end{cases}$$

Having more than 2 relations implying performing a cut at some point  $k$  and placing  $k$  on the left side,  $n-k$  on the right side. A factor of 2 indicates swapping the two sides.

This gives the closed form  $f(n) = 2^{n-1}C(n-1)$ .

### 3.0.2 Stars

Star queries have the constraint that one relation must be in the center; all the others can be ordered arbitrarily. This leads to the following formulas:

- Left-deep:  $2 \cdot (n-1)!$ , since there are  $n-1$  choices for a join partner and a factor of 2 for commutativity;
- Zig-zag:  $2 \cdot (n-1)! \cdot 2^{n-2}$ , in which the last factor represent the possibility to swap left and right for each subtree;
- Bushy trees: not possible since they require the first relation to be available.

### 3.0.3 Cliques

Cliques are a schema which do not care about cross products, since every relation is connected to the other, hence the number of possibilities is the same as the one obtained allowing cross products.

Still, complexity is very high and runtime is bad, although the worst case usually does not happen.

## 3.1 Greedy heuristics

Regardless of the methods and the inclusion of cross products, the search space is in general quite large when the number of relations is larger than 10, and polynomial time is hard to achieve. Some cost functions do not even have a proof of complexity.

Due to the size of the search space, greedy heuristics are ways to easily construct a potential tree in a fast time: they are most suitable for large queries, but often do not give the best result.

Algorithms assume no cross products within left-deep trees, and known cardinalities (or some other weight function).

### 3.1.1 GreedyJoinOrdering-1

This algorithm returns a set of ordered relations to be joined according to a cost function in a bushy tree structure.

The output is given starting from the minimum-weight relation, removes it and searches for the new minimum.

This method is simple, but not that good in practice: it assumed fixed weight (not depending on the size, for instance) and does not support optimization within intermediate results.

### 3.1.2 GreedyJoinOrdering-2

This variant also considers the previous set of relations, computing relative weights based on the existing tree.

In this case, however, the very first relation has a large relative weight and a major impact on the choices made afterwards.

### 3.1.3 GreedyJoinOrdering-3

To tackle this problem, a double loop is introduced in which not only the set of relations is scanned, but also every relation is tested as a starting one. Then, all computed results are compared and the minimum among them is returned.

This method is overall the best one and is implemented in some systems, but it is still not optimal.

### 3.1.4 Greedy Operator Ordering

Intermediate join trees must be combined to obtain larger trees: since some algorithms construct left-deep structures and others return bushy trees, those can be attached if the result is minimal: trees with the smallest weight are iteratively joined and then removed by the set of possibilities.

First, each combination of weight (for instance, product of cardinalities and selectivity) is calculated; then the minimum is chosen and the two relations are aggregated.

In case both relations are linked to the same (different) one in the query graph, the selectivity of the new edge consists in the product of previous ones.

This algorithm has two pitfalls: one is its computational complexity of  $O(n^3)$ , which however can be in the order of milliseconds with small input, and the other is its sub-optimal output in some cases.

In fact, it is not guaranteed that Greedy Operator Ordering finds the best bushy tree given a set of relation, even if it tends to have better results than the greedy ordering techniques.

## 3.2 IKKBZ

IKKBZ is a polynomial time algorithm for join ordering without cross products, producing left-deep trees under some assumptions: acyclic graphs, ASI cost functions and a fixed join technique.

The algorithm will ultimately compute a rank for each predicate, based on its selectivity, obtaining an optimal evaluation order. After the first steps, however, the remaining arguments are independent on the size of the relations.

It starts considering a cost function as a product of the form:

$$C(T_i \bowtie R_j) = |T_i| \cdot h_j(|R_j|)$$

Each relation  $R$  can have its own  $h$ , which has a set of parametrized cost functions. Cardinalities are also taken into account, defining  $n_i = |R_i|$  and  $h_i(n_i)$  as the cost per input tuple of a join.  $T_i$  is the left side of the computation.

A precedence graph helps identifying which relations should be joined first; it is represented as an oriented query graph, and constructed with root in  $R_k$  in the following way:

1. The root is fixed, removed by the set of possibilities and added to the tree;
2. As long as there are relations to be chosen,  $R_i \in V \setminus V_k^P$  is selected such that  $\exists R_j \in V_k^P : (R_j, R_i) \in E$ ;
3.  $R_i$  is added to  $V_k^P$  and an edge  $R_j \rightarrow R_i$  is created.

The algorithm constructs left-deep trees.

A sequence of nodes conforms to a precedence graph if two conditions are satisfied:

- For each position  $i$ , there exists one  $j$  coming before;
- There does not exist a position  $i$  and a  $j$  coming after with an edge from  $j$  to  $i$  (acyclic).

If there exists a path between two sets of relations  $R_i$  and  $R'_i$ , all relations from the first must be joined first; there is no join condition between any pair of relations aside from the one joining the two sets, hence the selectivity is 1.

Selectivity of the join:

$$s_i = \begin{cases} 1 & |R'_i| = 0 \\ \prod_{R_j \in R'_i} f_{i,j} & |R'_i| > 0 \end{cases}$$

In case of cycles, the selectivity cannot be uniquely determined: there could be two relations with different values. Therefore, first the precedence graph is fixed, and then selectivities are calculated.

If the query graph is a chain (total order), then the following properties hold:

$$n_{1,2,\dots,k} = \prod_{i=1}^k s_i n_i$$

$$C_H(G) = \sum_{i=2}^n [n_{1,2,\dots,i-1} h_i(n_i)] = \sum_{i=2}^n \left[ \left( \prod_{j=1}^i s_j n_j \right) h_i(n_i) \right]$$

The cardinality of the joins is equal to the product of each selectivity for each cardinality among all relations.

If the weight function is indeed the selectivity, then  $C_H \equiv C_{out}$ . The factor  $s_i n_i$  determines how much the input relation changes its cardinality before further joins, and can be increasing or decreasing.

The algorithm employs a recursive definition of the cost function:

$$\begin{cases} C_H(\epsilon) = 0 \\ C_H(R_i) = 0 & \text{the relation is the root} \\ C_H(R_i) = h_i(n_i) & \text{else} \\ C_H(S_1 S_2) = C_H(S_1) + T(S_1) \cdot C_H(S_2) \end{cases}$$

$$T(\epsilon) = 1 \quad \wedge \quad T(S) = \prod_{R_i \in S} s_i n_i$$

Last  $C_H$  definition corresponds to the cardinality performing all the relation in the sequence (the term  $T$ ) and multiplying for the cost of last element.

This allows to use ASI properties: in fact, these hold if and only if there exists a function  $T$  and a rank function defined as:

$$rank(S) = \frac{T(S) - 1}{C(S)}$$

The following must hold:

$$C(AUVB) \leq C(AVUB) \leftrightarrow rank(U) \leq rank(V)$$

In other words, ASI properties hold if and only if the relations can be ordered by rank. The cost function previously defined can be proven to respect these constraints.

It is possible that the rank contradicts the precedence graph: the following definition is introduced to counter this occurrence.

Let  $M = \{A_1, \dots, A_n\}$  be a set of sequences of nodes. Then  $M$  is called a module if, for all sequences  $B$  that do not overlap with the sequences in  $M$ , one of the following conditions holds:

- $B \rightarrow A_i, \forall A_i \in M$ ;
- $A_i \rightarrow B, \forall A_i \in M$ ;
- $B \nrightarrow A_i \wedge A_i \nrightarrow B, \forall A_i \in M$ .

If  $A \rightarrow B$  and  $rank(B) \leq rank(A)$ , then it is possible to find an optimal sequence among those in which  $B$  directly follows  $A$ .

The continued process of building a compound relation until no more contradictory sequences exist is called normalization; the opposite is denormalization.

### 3.2.1 The algorithm

IKKBZ works performing the following steps for each relation, considering it as a root node:

1. Calculates the precedence graph;
2. Executes the subprocedure of finding the subtree all of whose children are chains, and normalizing it, then merging the chains;
3. The relation returned by the previous step is added to the set.

The algorithm stops when all trees are single chains, then returns the minimum of the sets according to the cost function.

The subprocedure constructs a left-deep tree (chain) from the precedence graph, performing a normalization operation and merging based on the rank in ascending order. Normalization happens when there is a contradictory sequence, i. e. one chain has bigger rank than the following.

This works by taking  $r$  and  $c$  such that  $rank(r) > rank(c)$  and swapping them with a compound relation representing both. This allows to merge relations that would have been reordered if only considering the rank, obtaining the actual ascending order.

In case the graph contains cycles, it is possible to preprocess it with algorithms such as Minimum Spanning Tree to find a suitable representation to run IKKBZ.

### 3.3 Maximum Value Precedence

Maximum Value precedence is useful in those cases where IKKBZ fails, such as cyclic queries. It runs in polynomial time as well, employing a graph theoretic approach to calculate in how many ways joins can be scheduled.

The algorithm requires a join graph, similar to the ones discussed before, which will be modified later, hence the definition must be extended: edges are assigned an order, and predicates are used to identify sets.


If there exists an edge between the predicates  $p_1$  and  $p_2$ , and they belong to the same relation, two directed edges are added in the graph.

If  $p_1$  is connected to  $p_2$  and  $p_2$  is connected to  $p_3$ , there will be a so called virtual edge connecting  $p_1$  and  $p_3$ : eventually, all nodes will be connected with either physical or virtual edges, forming a clique.

For instance, every spanning tree in the directed join graph leads to a join tree: first of all edges are added in both direction, and then one of them gets removed, giving a directed acyclic graph which can easily give a join order (without distinguishing left and right).

MVP has the advantage of also producing bushy trees, however it does not guarantee optimality of the spanning tree; furthermore, the output might not correspond to an effective join tree.

Despite the incorrect representation, it is simple to fix this kind of errors, but there is no specific way to produce the best join tree.



To remedy the uncertainty, some additional rules are introduced to identify an effective spanning tree. The following conditions must be satisfied:

1.  $T$  must be binary (no nodes can have more than two children);
2. For all inner connected nodes  $(u, v)$ ,  $R(T(u)) \cap R(v) \neq \emptyset$  (they must have predicates in common);
3. For all  $(u_1, v)$ ,  $(u_2, v)$  one of the following holds:
  - (a)  $((R(T(u_1)) \cap R(v)) \cap ((R(T(u_2)) \cap R(v))) = \emptyset$  (they have a different relation in common);
  - (b)  $(R(T(u_1)) = R(v)) \wedge (R(T(u_2)) = R(v))$  (the relations in common are the same between pairs).

An effective spanning tree corresponds to a valid join tree, despite the definition not being intuitive. Given this, the rest of the assumptions is simple: if every predicate involves two relations and an equi-join condition, then  $R(u) \cap R(v)$  contains only a single relation.

Let  $v$  be that relation, then  $R_i \bowtie_v R_j$  is abbreviated by  $\bowtie_v$ .

The next step is adding weight to the edges, obtaining a weighted directed join graph. Each weight

is calculated with the formula:

$$w_{u,v} = \frac{|\bowtie_u|}{R(u) \cap R(v)}$$

This means that each edge has a weight depending on the relation they have in common and the join cardinality. The intuition implies that a join is executed before another, the given relation becomes available and the edge weight can explain how the cardinality changes, i. e. how many tuples are generated.

This value can be bigger or smaller than 1, and can make the total cost bigger or smaller (the input size changes by a factor of  $w_{u,v}$ ). For virtual edges, the weight is 1 by default.

It is relevant to notice that the weight function used by MVP is the same one as  $s_i$  in IKKBZ. However, it can also be chosen arbitrarily.

Of course, a weight smaller than 1 reduces the cost of following join operations. Furthermore, weights change over time depending on a partial spanning tree:

$$w(p_{i,j}, S) = \frac{|\bowtie_{p_{i,j}}^S|}{|R_i \bowtie_{p_{i,j}} R_j|}$$

$\bowtie_{p_{i,j}}^S$  is the result of the join after all joins preceding  $p_{i,j}$  in  $S$  have been executed. If the spanning tree is empty, the cost is merely equal to the cost of a simple join.

### 3.3.1 The algorithm

The algorithm works in two phases:

1. Taking the edges with weight smaller than 1, trying to reduce the work for latter operators as soon as possible;
2. Adding the remaining edges, potentially causing an increase of the load, yet as late as possible.

MVP takes as input a weighted directed join graph and builds two priority queues, one with largest weights (phase 1) and one with smallest (phase 2).

The working graph is initially just a set of predicates with no edges, and then the two phases are ran.

The first phase modifies the state of a working tree, taking the head of the priority queue (the most expensive) and finding the joins which could make it cheaper, adding nodes with smaller weight while still keeping the graph acyclic.

If there is no such edge, the element is swapped from the first queue to the second; else, the edge to be added to the working tree will be the one which maximizes the difference within costs (minimizes the new cost). Weights are then recomputed.

Phase 2 is called when the second queue is non-empty, again considering edges which respect the acyclic property. The procedure tries to minimize the additional cost caused by adding joins.

To effectively modify the working tree, an update function changes the state by performing unions of sets with edges and removal from the graph yet to consider. If there are two incoming physical edges, they are replaced with virtual ones.

This also ensures the binary property, removing all cases in which a node could have two parents and handling eventual duplicates. However, there is still no guarantee of optimality.




### 3.4 Dynamic Programming

Dynamic programming approaches can be useful to obtain more insights on possible orders with cost function. Since this kind of algorithm is a macro class, they can be defined on any input and give any output (bushy, left-deep).

These work thanks to an optimality principle: if an option is cheaper than another, the latter can be discarded and the best solution is found only considering the set of further possibilities generated by the first.

Formally, let  $T$  be an optimal join tree for relations  $R_1, \dots, R_n$ . Then, every subtree  $S$  of  $T$  is an optimal join tree for the relations contained in it.

Despite some hypothetical concerns of sub-optimality and the presence of physical properties which may alter the result, in practice this property holds.



The strategy works starting with a single relation and generating larger trees with a bottom-up strategy, reusing previous intermediate results.

A possible dynamic programming algorithm calculates the cost functions when joining either on the left or right side (bushy usually works better) for each applicable join implementation.

The outcome will be a list of pairs with their cost, of which the minimum is chosen and propagated through following iterations.

The search space gets therefore reduced whenever an option is discarded, hence the number of combinations is never exponential.

In the case of linear trees, there is a basic strategy finding the optimal  $T$  by joining all optimal  $T'$  with  $T \setminus T'$ , with  $|T| = |T'| + 1$ .

The most common algorithm constructs the optimal left-deep tree from an empty table mapping the set of relations ( $2^R$ ) to the join tree.

For each relation, the optimal tree is built, starting from size 1 and increasing the subsets by adding those of smaller size having the optimal solution of the subproblem.